

TEAM ID : PNT2022TMID32010

PROJECT TITLE : Estimate The Crop Yield Using Data Analytics

ASSIGNMENT DATE : 31.10.22

STUDENT NAME : MANOJ J

STUDENT ROLL NUMBER : 731619104029

1.Download the dataset

In [265]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

2.LOAD THE DATASET

In [266]:

```
df = pd.read_csv('abalone.csv')
```

In [267]:

```
df.head
```

Out[267]:

```
<bound method NDFrame.head of          Sex  Length  Diameter  Height  Whole weight  Shucked w
eight  \
0      M    0.455    0.365    0.095          0.5140          0.2245
1      M    0.350    0.265    0.090          0.2255          0.0995
2      F    0.530    0.420    0.135          0.6770          0.2565
3      M    0.440    0.365    0.125          0.5160          0.2155
4      I    0.330    0.255    0.080          0.2050          0.0895
...
4172   F    0.565    0.450    0.165          0.8870          0.3700
4173   M    0.590    0.440    0.135          0.9660          0.4390
4174   M    0.600    0.475    0.205          1.1760          0.5255
4175   F    0.625    0.485    0.150          1.0945          0.5310
4176   M    0.710    0.555    0.195          1.9485          0.9455

      Viscera weight  Shell weight  Rings
0              0.1010        0.1500      15
1              0.0485        0.0700       7
2              0.1415        0.2100       9
3              0.1140        0.1550      10
4              0.0395        0.0550       7
...
4172            0.2390        0.2490      11
4173            0.2145        0.2605      10
4174            0.2875        0.3080       9
4175            0.2610        0.2960      10
4176            0.3765        0.4950      12
```

```
[4177 rows x 9 columns]>
```

In [268]:

```
Age=1.5+df.Rings
```

```
df["Age"]=Age
df=df.rename(columns = {'whole weight':'whole_weight','Shucked weight':'Shucked_weight',
'Viscera weight':'Viscera_weight','Shell weight':'Shell_weight'})
df=df.drop(columns=["Rings"],axis=1)
df.head()
```

Out[268]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 0 | M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 16.5 |
| 1 | M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 8.5 |
| 2 | F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 10.5 |
| 3 | M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 11.5 |
| 4 | I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 8.5 |

In [269]:

```
df.tail()
```

Out[269]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|------|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 4172 | F | 0.565 | 0.450 | 0.165 | 0.8870 | 0.3700 | 0.2390 | 0.2490 | 12.5 |
| 4173 | M | 0.590 | 0.440 | 0.135 | 0.9660 | 0.4390 | 0.2145 | 0.2605 | 11.5 |
| 4174 | M | 0.600 | 0.475 | 0.205 | 1.1760 | 0.5255 | 0.2875 | 0.3080 | 10.5 |
| 4175 | F | 0.625 | 0.485 | 0.150 | 1.0945 | 0.5310 | 0.2610 | 0.2960 | 11.5 |
| 4176 | M | 0.710 | 0.555 | 0.195 | 1.9485 | 0.9455 | 0.3765 | 0.4950 | 13.5 |

3. Perform Below Visualizations


Univariate Analysis

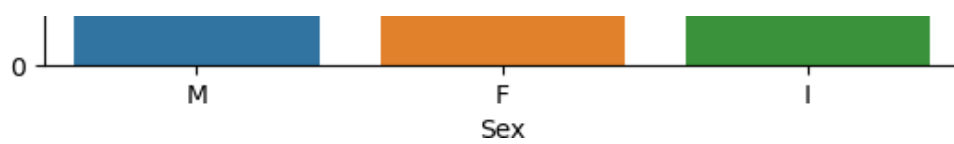
In [270]:

```
sns.countplot(x='Sex',data=df)
```

Out[270]:

<AxesSubplot:xlabel='Sex', ylabel='count'>



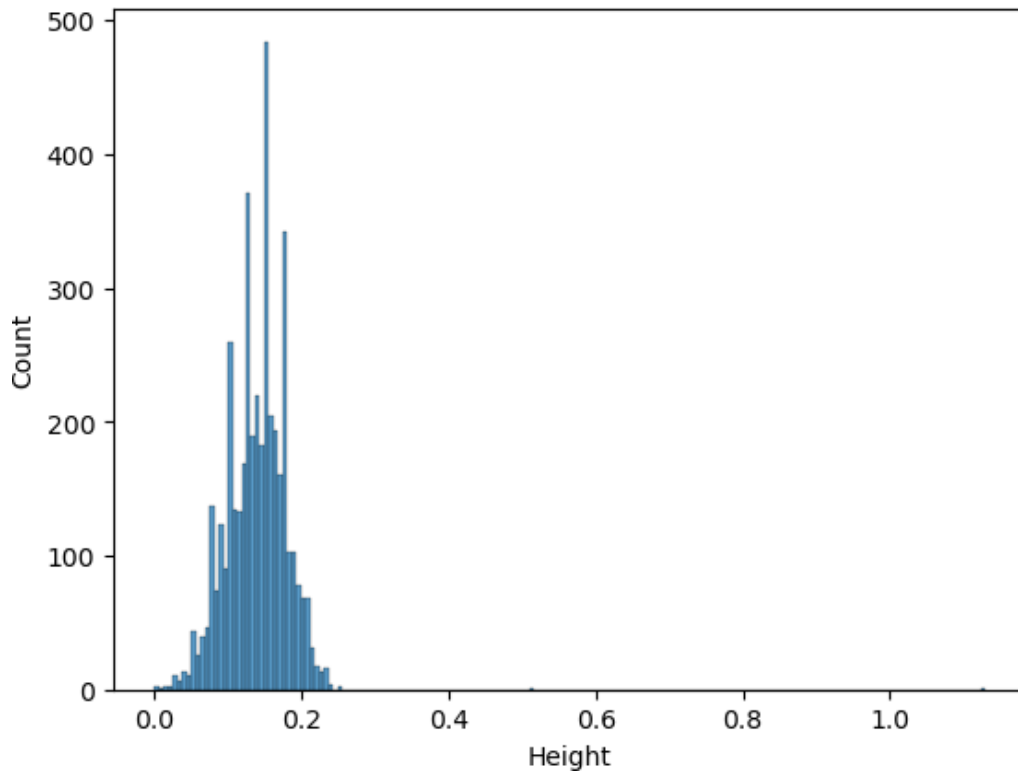


In [271]:

```
sns.histplot(df["Height"])
```

Out[271]:

<AxesSubplot:xlabel='Height', ylabel='Count'>

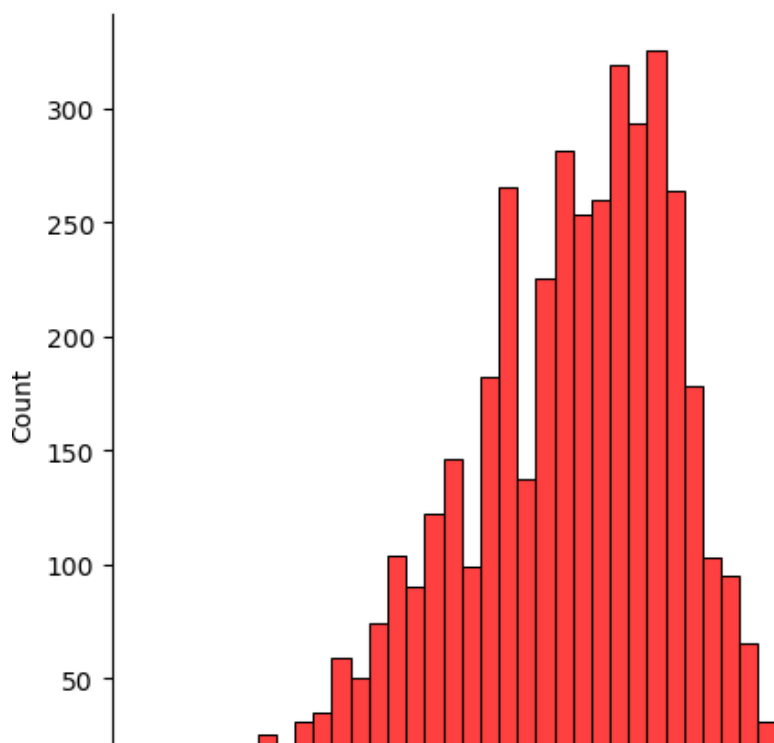


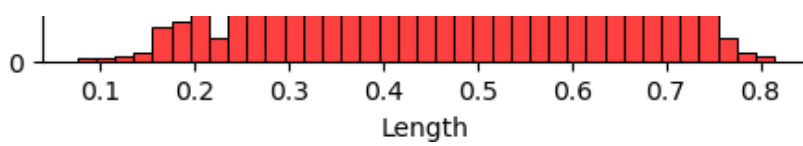
In [272]:

```
sns.displot(df["Length"], color='red')
```

Out[272]:

<seaborn.axisgrid.FacetGrid at 0x1af5e2f7820>

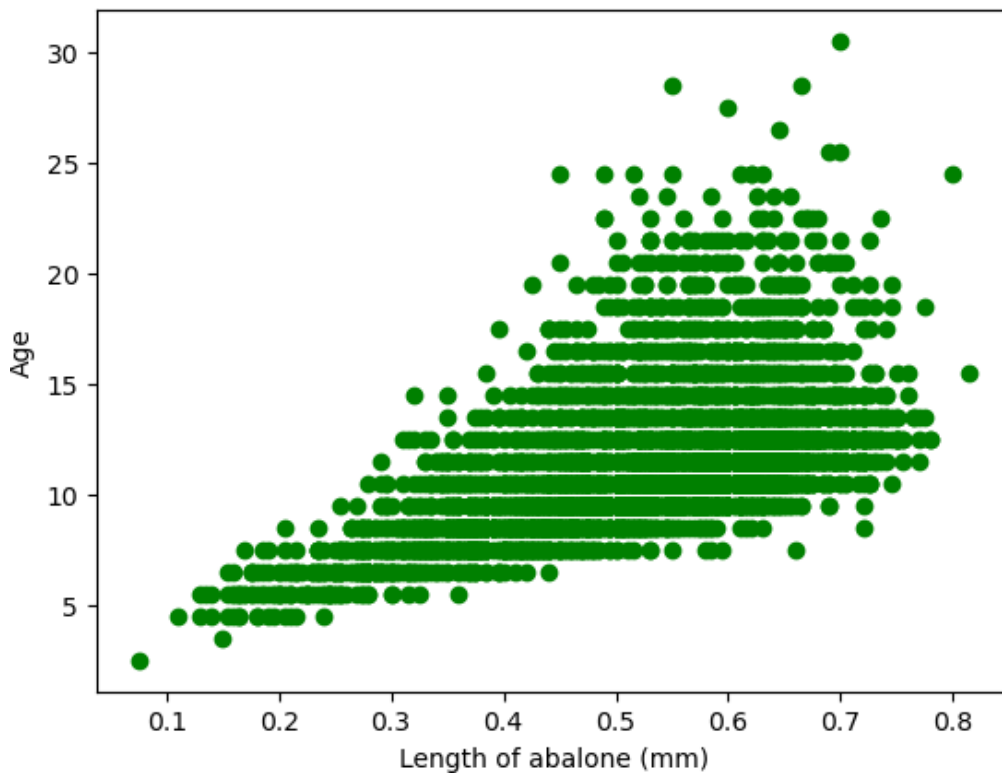




Bi-Variate Analysis

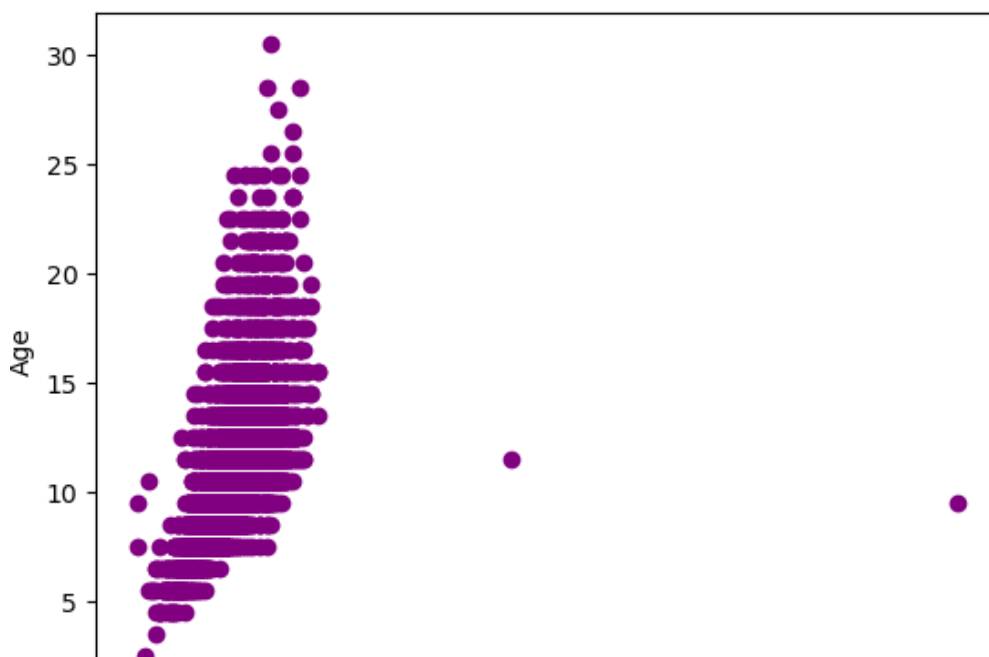
In [273]:

```
plt.scatter(df['Length'],df['Age'],c='green')
plt.xlabel('Length of abalone (mm)')
plt.ylabel('Age')
plt.show()
```



In [274]:


```
plt.scatter(df['Height'],df['Age'],c='purple')
plt.xlabel('Height of abalone (mm)')
plt.ylabel('Age')
plt.show()
```



0.0 0.2 0.4 0.6 0.8 1.0
Height of abalone (mm)

In [275]:

```
plt.scatter(df['Diameter'],df['Age'],c='black')  
plt.xlabel('Height of abalone (mm)')  
plt.ylabel('Age')  
plt.show()
```



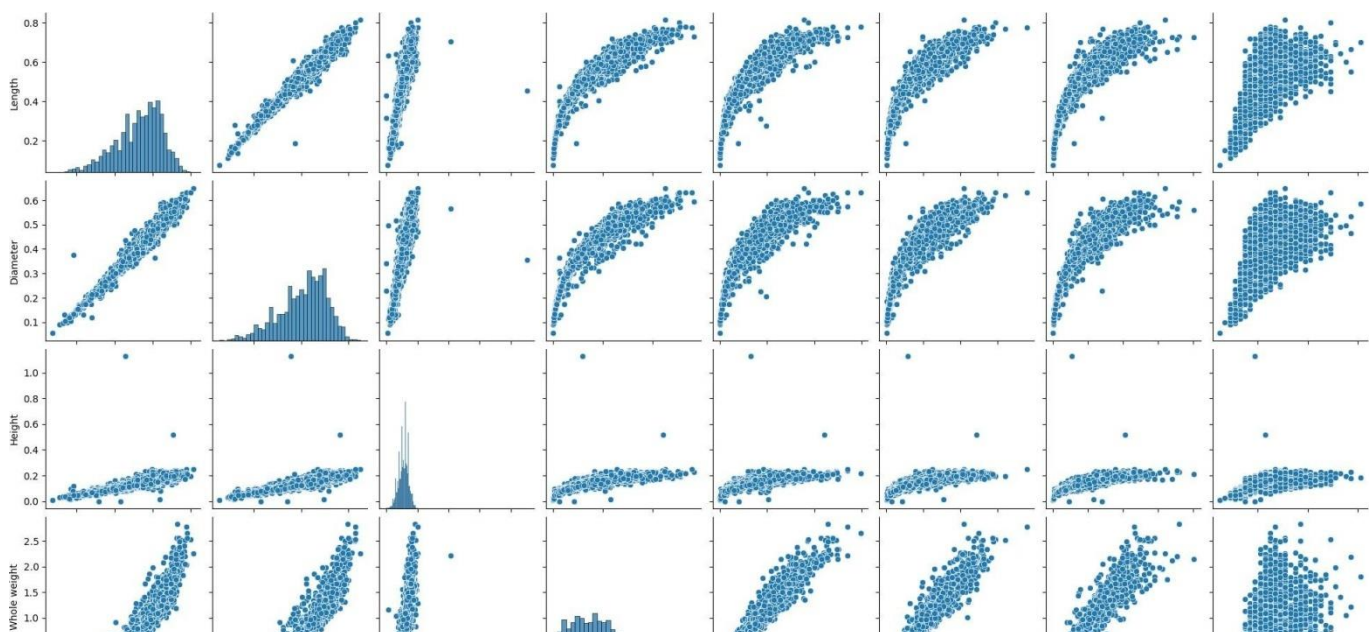
Multi-Variate Analysis

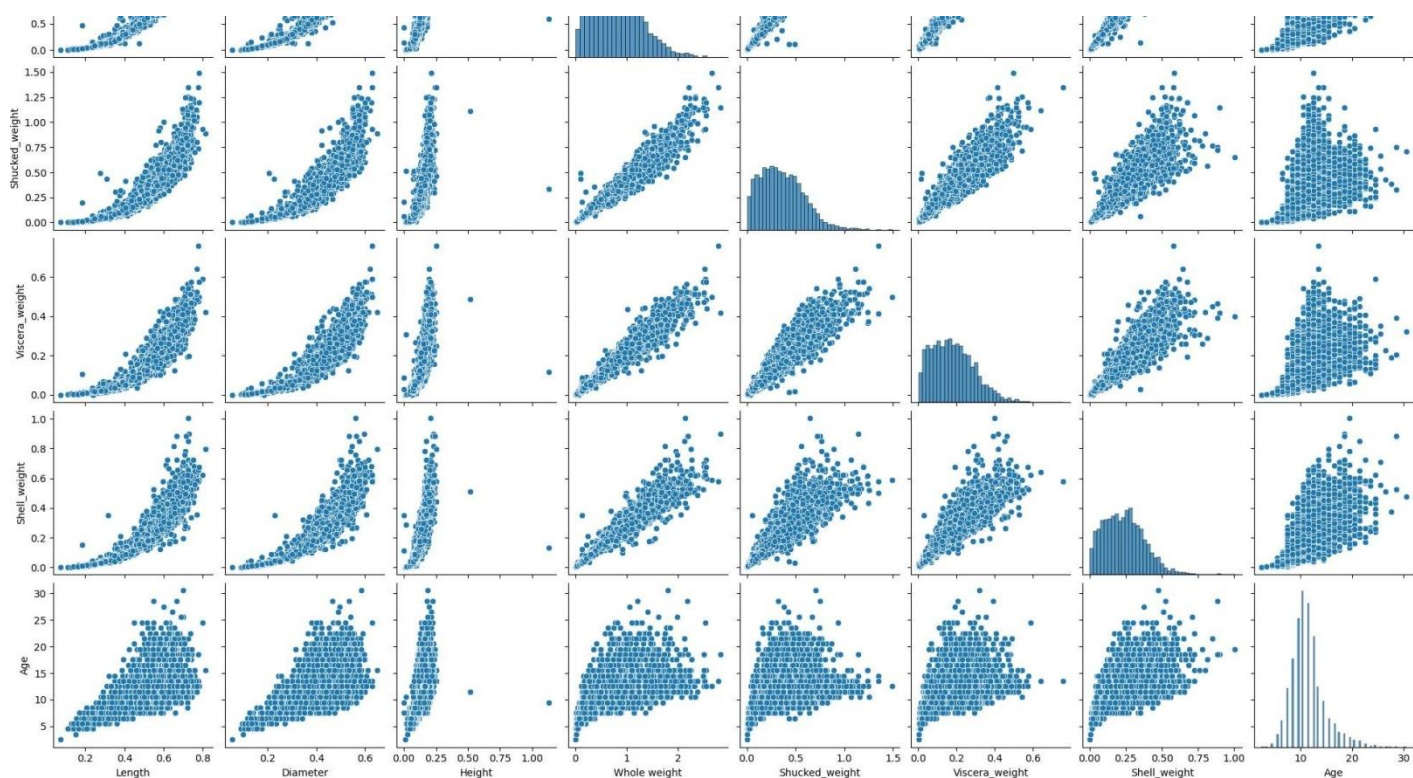
In [276]:

```
numerical_features = df.select_dtypes(include = [np.number]).columns  
sns.pairplot(df[numerical_features])
```

Out[276]:

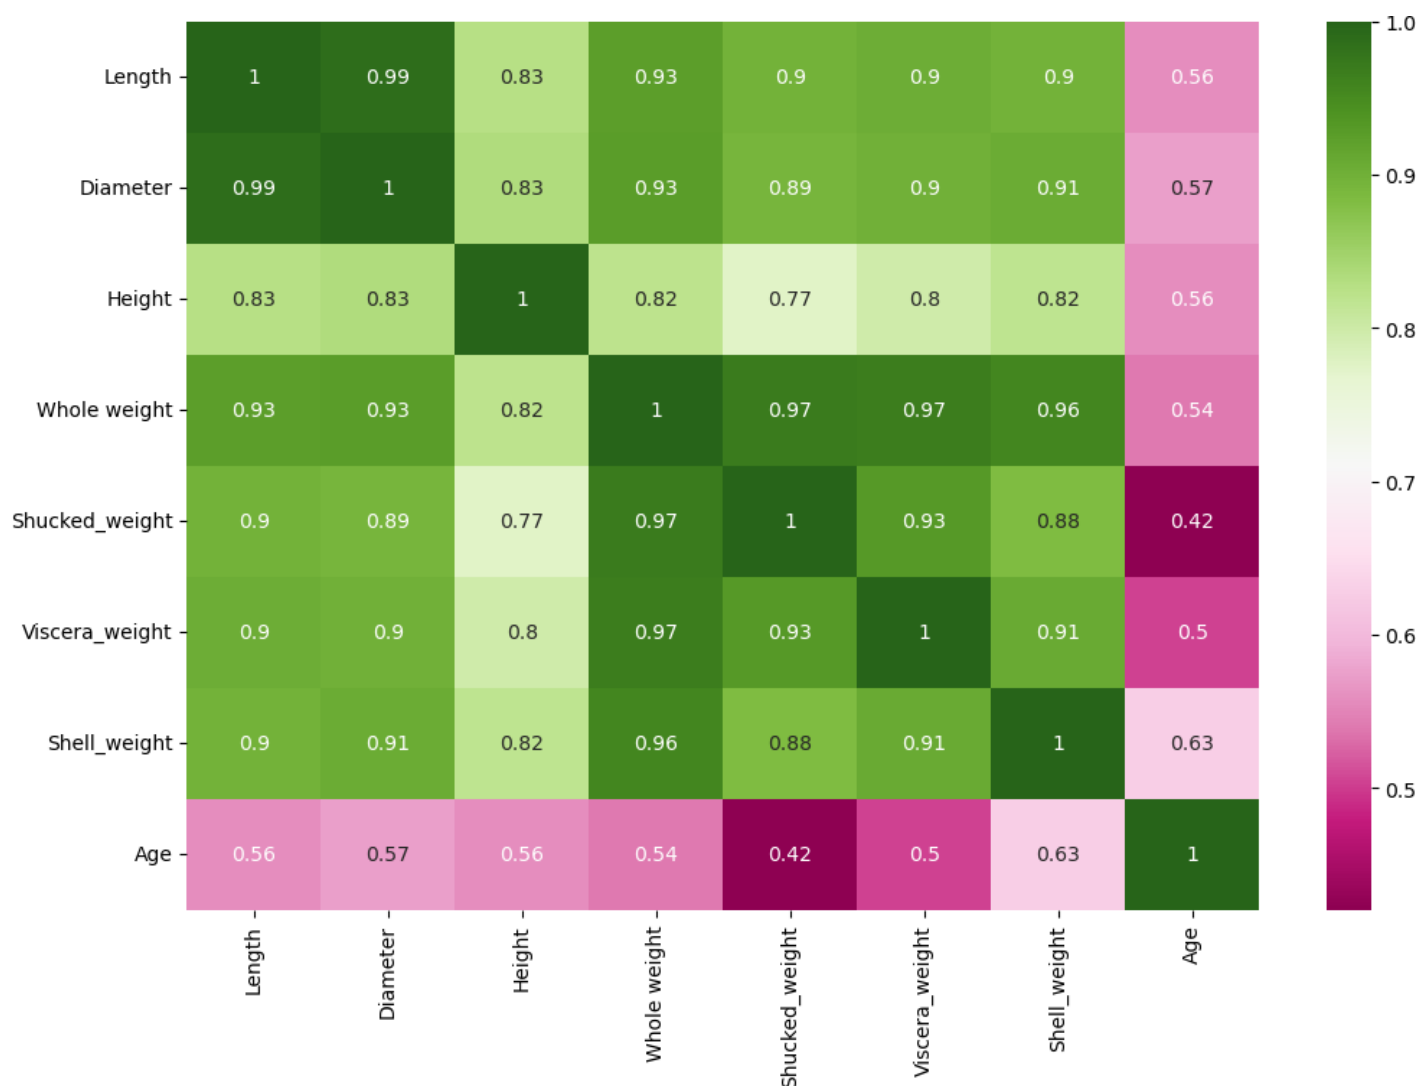
<seaborn.axisgrid.PairGrid at 0x1af61732d00>





In [277]:

```
plt.figure(figsize=(12,8));
sns.heatmap(df.corr(),cmap='PiYG',annot=True);
```



4. Perform descriptive statistics on the dataset

In [278]:

```
df.describe()
```

Out[278]:

| | Length | Diameter | Height | Whole weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|-------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|-------------|
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 |
| mean | 0.523992 | 0.407881 | 0.139516 | 0.828742 | 0.359367 | 0.180594 | 0.238831 | 11.433684 |
| std | 0.120093 | 0.099240 | 0.041827 | 0.490389 | 0.221963 | 0.109614 | 0.139203 | 3.224169 |
| min | 0.075000 | 0.055000 | 0.000000 | 0.002000 | 0.001000 | 0.000500 | 0.001500 | 2.500000 |
| 25% | 0.450000 | 0.350000 | 0.115000 | 0.441500 | 0.186000 | 0.093500 | 0.130000 | 9.500000 |
| 50% | 0.545000 | 0.425000 | 0.140000 | 0.799500 | 0.336000 | 0.171000 | 0.234000 | 10.500000 |
| 75% | 0.615000 | 0.480000 | 0.165000 | 1.153000 | 0.502000 | 0.253000 | 0.329000 | 12.500000 |
| max | 0.815000 | 0.650000 | 1.130000 | 2.825500 | 1.488000 | 0.760000 | 1.005000 | 30.500000 |

5.Check for Missing values and deal with them

In [279]:

```
df.isnull().sum()
```

Out[279]:

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked_weight 0
Viscera_weight 0
Shell_weight 0
Age          0
dtype: int64
```

6.Find the outliers and replace them outliers


In [280]:

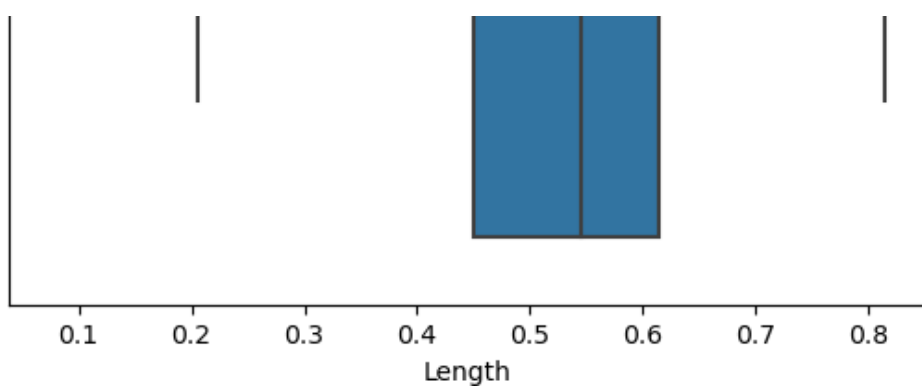
```
sns.boxplot(df['Length'])
```

```
C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
```

Out[280]:

```
<AxesSubplot:xlabel='Length'>
```





In [281]:

```
q1 = df['Length'].quantile(0.25)
q2 = df['Length'].quantile(0.75)
iqr = q2-q1
q1, q2,iqr
```

Out[281]:

```
(0.45, 0.615, 0.16499999999999998)
```

In [282]:

```
upper_limit = q2+(1.5*iqr)
lower_limit = q1-(1.5*iqr)
lower_limit, upper_limit
```

Out[282]:

```
(0.20250000000000004, 0.8624999999999999)
```

In [283]:

```
new_df = df.loc[(df['Length'] <= upper_limit) & (df['Length'] >= lower_limit)]
print('before removing outliers:', len(df))
print('after removing outliers:', len(new_df))
print('outliers:', len(df)-len(new_df))
```

```
before removing outliers: 4177
after removing outliers: 4128
outliers: 49
```

In [284]:

```
new_df = df.copy()
new_df.loc[(new_df['Length']>upper_limit), 'Length'] = upper_limit
new_df.loc[(new_df['Length']<lower_limit), 'Length'] = lower_limit
```

In [285]:


```
sns.boxplot(new_df['Length'])
```

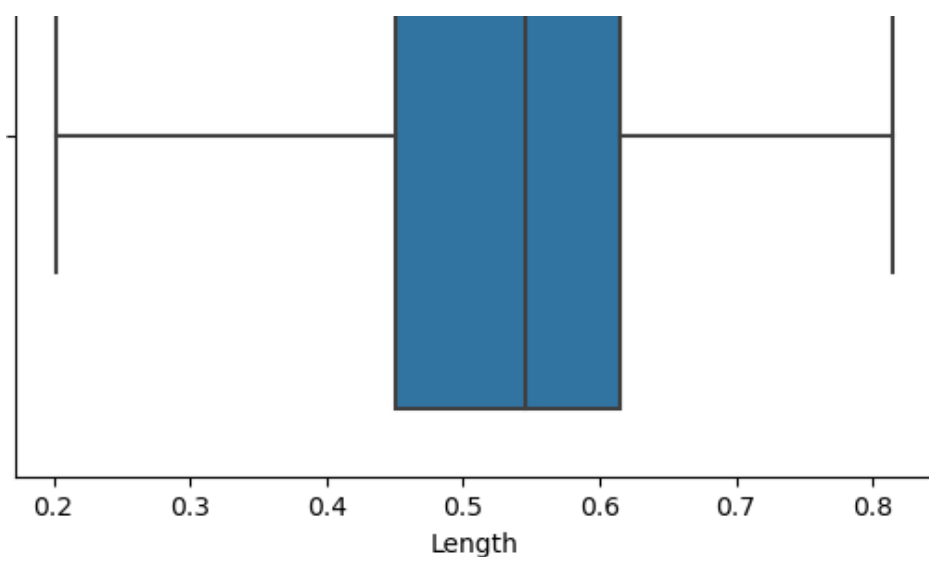
C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[285]:

```
<AxesSubplot:xlabel='Length'>
```





In [286]:


```
sns.boxplot(df['Diameter'])
```

C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[286]:

```
<AxesSubplot:xlabel='Diameter'>
```



In [287]:

```
q1 = df['Diameter'].quantile(0.25)
q2 = df['Diameter'].quantile(0.75)
iqr = q2-q1
q1, q2, iqr
```

Out[287]:

```
(0.35, 0.48, 0.13)
```

In [288]:

```
upper_limit = q2 + (1.5*iqr)
lower_limit = q1 - (1.5*iqr)
lower_limit, upper_limit
```

Out[288]:

```
(0.15499999999999997, 0.675)
```

In [289]:

```
new_df = df.loc[(df['Diameter'] <= upper_limit) & (df['Diameter'] >= lower_limit)]
print('before removing outliers :', len(df))
print('after removing outliers :', len(new_df))
print('outliers :', len(df)-len(new_df))
```

```
before removing outliers : 4177
after removing outliers : 4118
outliers : 59
```

In [290]:

```
new_df = df.copy()
new_df.loc[(new_df['Diameter']>upper_limit), 'Diameter'] = upper_limit
new_df.loc[(new_df['Diameter']<lower_limit), 'Diameter'] = lower_limit
```

In [291]:


```
sns.boxplot(new_df['Diameter'])
```

C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[291]:

```
<AxesSubplot:xlabel='Diameter'>
```



In [292]:


```
sns.boxplot(df['Height'])
```

C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will res

```
ult in an error or misinterpretation.  
warnings.warn(
```

```
Out[292]:
```

```
<AxesSubplot:xlabel='Height'>
```



```
In [293]:
```

```
q1 = df['Height'].quantile(0.25)  
q2 = df['Height'].quantile(0.75)  
iqr = q2-q1  
q1, q2, iqr
```

```
Out[293]:
```

```
(0.115, 0.165, 0.05)
```

```
In [294]:
```

```
upper_limit = q2 + (1.5*iqr)  
lower_limit = q1 - (1.5*iqr)  
lower_limit, upper_limit
```

```
Out[294]:
```

```
(0.039999999999999994, 0.24000000000000002)
```

```
In [295]:
```

```
new_df = df.loc[(df['Height'] <= upper_limit) & (df['Height'] >= lower_limit)]  
print('before removing outliers :', len(df))  
print('after removing outliers :', len(new_df))  
print('outliers :', len(df)-len(new_df))
```

```
before removing outliers : 4177  
after removing outliers : 4148  
outliers : 29
```

```
In [296]:
```

```
new_df = df.copy()  
new_df.loc[(new_df['Height']>upper_limit), 'Height'] = upper_limit  
new_df.loc[(new_df['Height']<lower_limit), 'Height'] = lower_limit
```

```
In [297]:
```


```
sns.boxplot(new_df['Height'])
```

C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[297]:

<AxesSubplot:xlabel='Height'>



In [298]:


```
sns.boxplot(new_df['Height'])
```

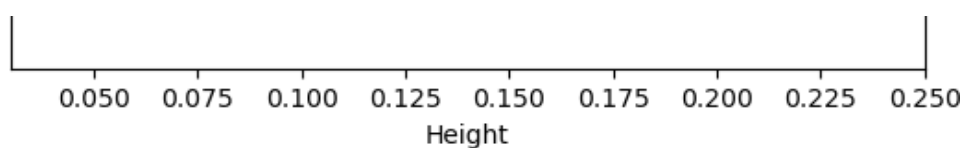
C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[298]:

<AxesSubplot:xlabel='Height'>





In [299]:

```
q1 = df['Whole weight'].quantile(0.25)
q2 = df['Whole weight'].quantile(0.75)
iqr = q2-q1
q1, q2, iqr
```

Out[299]:

```
(0.4415, 1.153, 0.7115)
```

In [300]:

```
upper_limit = q2 + (1.5*iqr)
lower_limit = q1 - (1.5*iqr)
lower_limit, upper_limit
```

Out[300]:

```
(-0.62575, 2.22025)
```

In [301]:

```
new_df = df.loc[(df['Whole weight'] <= upper_limit) & (df['Whole weight'] >= lower_limit)
]
print('before removing outliers :', len(df))
print('after removing outliers :', len(new_df))
print('outliers :', len(df)-len(new_df))
```

```
before removing outliers : 4177
after removing outliers : 4147
outliers : 30
```

In [302]:

```
new_df = df.copy()
new_df.loc[(new_df['Whole weight']>upper_limit), 'Whole weight'] = upper_limit
new_df.loc[(new_df['Whole weight']<lower_limit), 'Whole weight'] = lower_limit
```

In [303]:

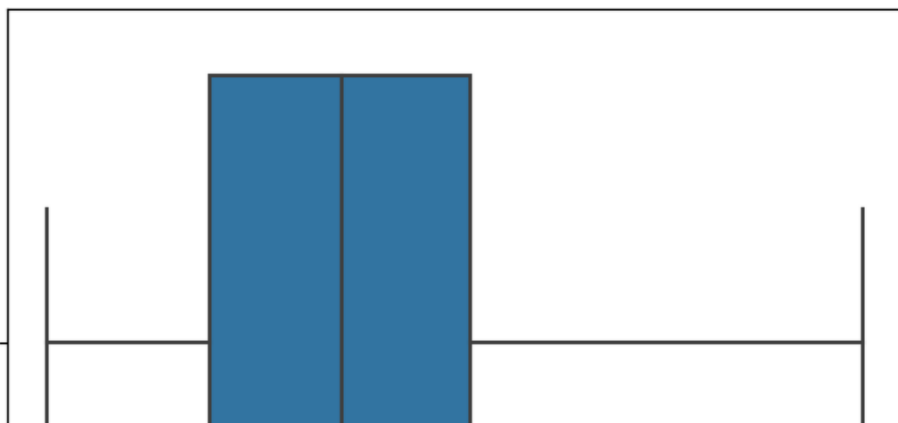
```
sns.boxplot(new_df['Whole weight'])
```

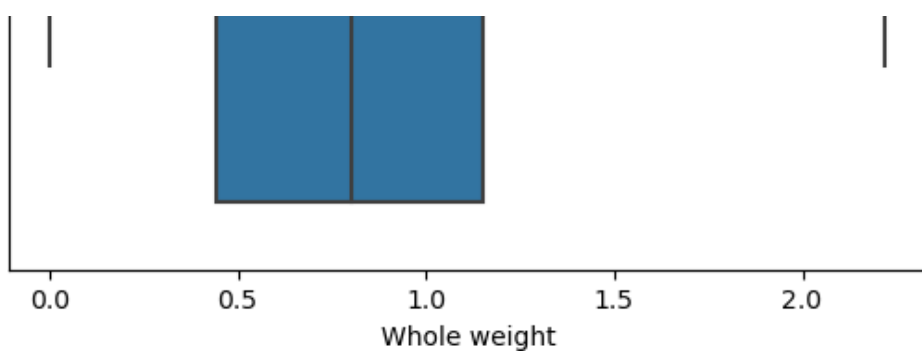
C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[303]:

```
<AxesSubplot:xlabel='Whole weight'>
```





In [304]:

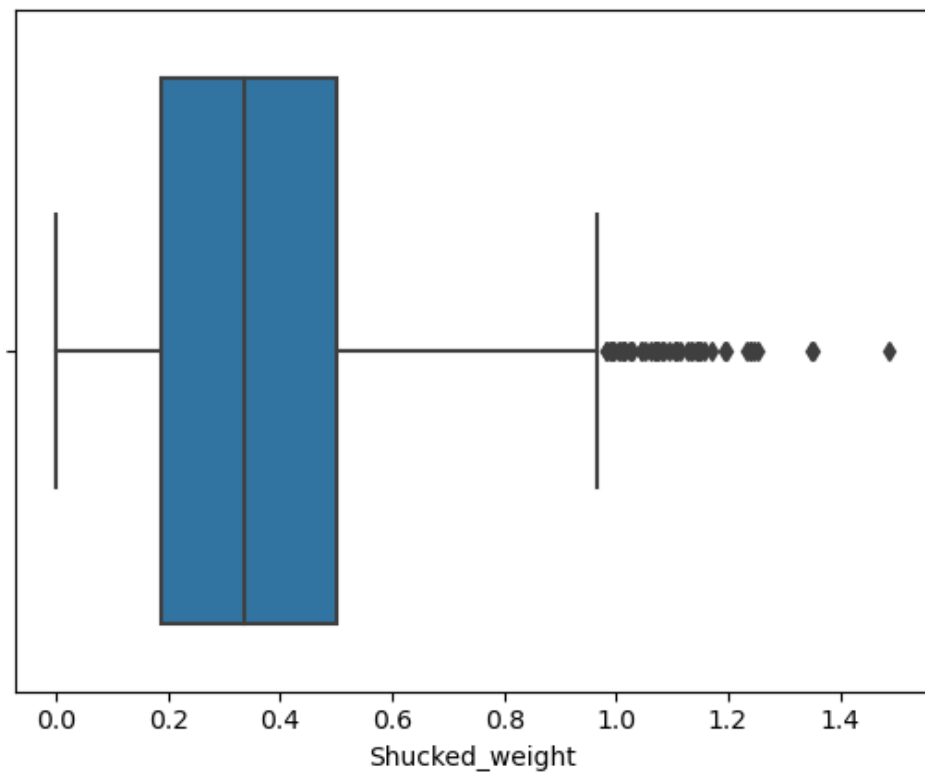
```
sns.boxplot(df['Shucked_weight'])
```

C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[304]:

```
<AxesSubplot:xlabel='Shucked_weight'>
```



In [305]:

```
q1 = df['Shucked_weight'].quantile(0.25)
q2 = df['Shucked_weight'].quantile(0.75)
iqr = q2-q1
q1, q2, iqr
```

Out[305]:

```
(0.186, 0.502, 0.316)
```

In [306]:

```
upper_limit = q2 + (1.5*iqr)
lower_limit = q1 - (1.5*iqr)
lower_limit, upper_limit
```

Out[306]:

```
(-0.288, 0.976)
```

In [307]:

```
new_df = df.loc[(df['Shucked_weight'] <= upper_limit) & (df['Shucked_weight'] >= lower_limit)]
print('before removing outliers :', len(df))
print('after removing outliers :', len(new_df))
print('outliers :', len(df)-len(new_df))
```

```
before removing outliers : 4177
after removing outliers : 4129
outliers : 48
```

In [308]:

```
new_df = df.copy()
new_df.loc[(new_df['Shucked_weight']>upper_limit), 'Shucked_weight'] = upper_limit
new_df.loc[(new_df['Shucked_weight']<lower_limit), 'Shucked_weight'] = lower_limit
```

In [309]:


```
sns.boxplot(new_df['Shucked_weight'])
```

C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[309]:

<AxesSubplot:xlabel='Shucked_weight'>



In [310]:

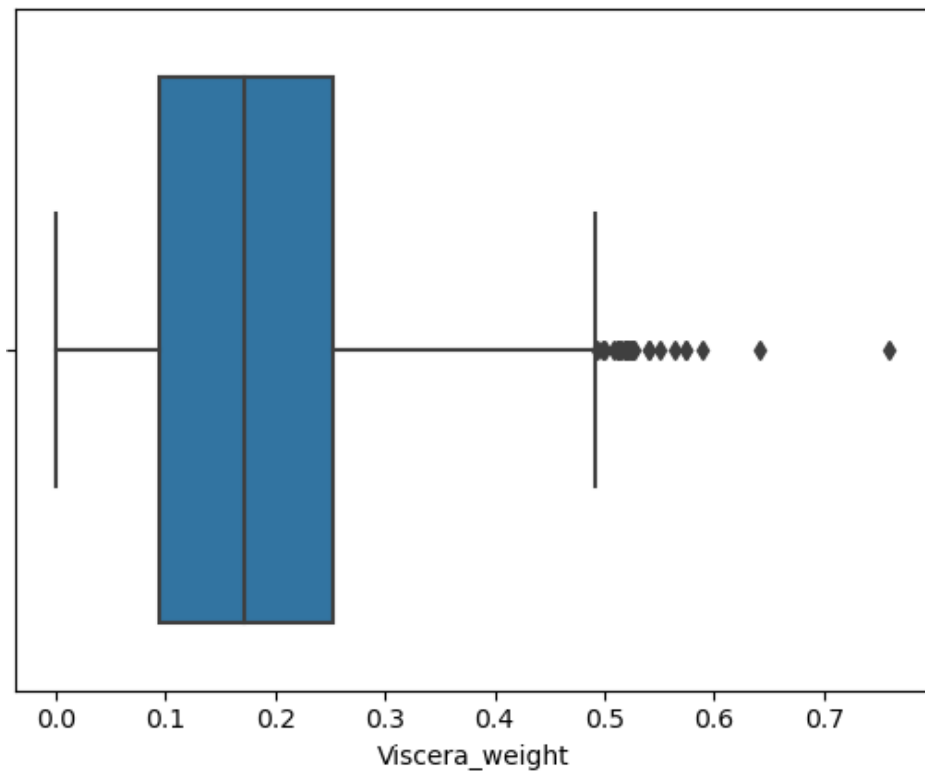
```
sns.boxplot(df['Viscera_weight'])
```

C:\Users\ELCOT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[310]:

<AxesSubplot:xlabel='Viscera_weight'>



7. Check for Categorical columns and perform encoding

In [311]:

```
df['Sex'].replace({'M':1, 'F':0, 'I':2}, inplace=True)
df
```

Out[311]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|------|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 0 | 1 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.1500 | 16.5 |
| 1 | 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.0700 | 8.5 |
| 2 | 0 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.2100 | 10.5 |
| 3 | 1 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.1550 | 11.5 |
| 4 | 2 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.0550 | 8.5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4172 | 0 | 0.565 | 0.450 | 0.165 | 0.8870 | 0.3700 | 0.2390 | 0.2490 | 12.5 |
| 4173 | 1 | 0.590 | 0.440 | 0.135 | 0.9660 | 0.4390 | 0.2145 | 0.2605 | 11.5 |
| 4174 | 1 | 0.600 | 0.475 | 0.205 | 1.1760 | 0.5255 | 0.2875 | 0.3080 | 10.5 |
| 4175 | 0 | 0.625 | 0.485 | 0.150 | 1.0945 | 0.5310 | 0.2610 | 0.2960 | 11.5 |
| 4176 | 1 | 0.710 | 0.555 | 0.195 | 1.9485 | 0.9455 | 0.3765 | 0.4950 | 13.5 |

4177 rows × 9 columns

In [312]:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
```

In [313]:

```
label_encoder = LabelEncoder()
df['Sex'] = label_encoder.fit_transform(df['Sex'])
df
```

Out[313]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|------|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 0 | 1 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.1500 | 16.5 |
| 1 | 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.0700 | 8.5 |
| 2 | 0 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.2100 | 10.5 |
| 3 | 1 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.1550 | 11.5 |
| 4 | 2 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.0550 | 8.5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4172 | 0 | 0.565 | 0.450 | 0.165 | 0.8870 | 0.3700 | 0.2390 | 0.2490 | 12.5 |
| 4173 | 1 | 0.590 | 0.440 | 0.135 | 0.9660 | 0.4390 | 0.2145 | 0.2605 | 11.5 |
| 4174 | 1 | 0.600 | 0.475 | 0.205 | 1.1760 | 0.5255 | 0.2875 | 0.3080 | 10.5 |
| 4175 | 0 | 0.625 | 0.485 | 0.150 | 1.0945 | 0.5310 | 0.2610 | 0.2960 | 11.5 |
| 4176 | 1 | 0.710 | 0.555 | 0.195 | 1.9485 | 0.9455 | 0.3765 | 0.4950 | 13.5 |

4177 rows × 9 columns

In [314]:

```
enc = OneHotEncoder(drop='first')

enc_df = pd.DataFrame(enc.fit_transform(df[['Sex']]).toarray())

df = df.join(enc_df)

df.head()
```

Out[314]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked_weight | Viscera_weight | Shell_weight | Age | 0 | 1 |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------|-----|-----|
| 0 | 1 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 16.5 | 1.0 | 0.0 |
| 1 | 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 8.5 | 1.0 | 0.0 |
| 2 | 0 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 10.5 | 0.0 | 0.0 |
| 3 | 1 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 11.5 | 1.0 | 0.0 |
| 4 | 2 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 8.5 | 0.0 | 1.0 |

8.Split the data into dependent and independent variables

In [315]:

```
x.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 11 columns):
#   Column              Non-Null Count  Dtype
----  -
0    Length              4177 non-null  float64
1    Diameter            4177 non-null  float64
2    Height              4177 non-null  float64
3    Whole weight        4177 non-null  float64
4    Shucked_weight      4177 non-null  float64
5    Viscera_weight      4177 non-null  float64
6    Shell_weight        4177 non-null  float64
7    Age                 4177 non-null  float64
8    Sex_F               4177 non-null  uint8
9    Sex_I               4177 non-null  uint8
10   Sex_M               4177 non-null  uint8
dtypes: float64(8), uint8(3)
memory usage: 273.4 KB
```

In [316]:

```
X = x.drop(['Age'], axis = 1)
```

In [317]:

```
X.head(2)
```

Out[317]:

| | Length | Diameter | Height | Whole weight | Shucked_weight | Viscera_weight | Shell_weight | Sex_F | Sex_I | Sex_M |
|---|--------|----------|--------|--------------|----------------|----------------|--------------|-------|-------|-------|
| 0 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.15 | 0 | 0 | 1 |
| 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.07 | 0 | 0 | 1 |

In [318]:

```
y = x['Age']
```

In [319]:

```
y.head(2)
```

Out[319]:

```
0    16.5
1     8.5
Name: Age, dtype: float64
```

9. Scale the independent variables

In [320]:

```
scale = StandardScaler()
scaledX = scale.fit_transform(x)

print(scaledX)
```

```
[[-0.57455813 -0.43214879 -1.06442415 ... -0.67483383 -0.68801788
  1.31667716]
 [-1.44898585 -1.439929   -1.18397831 ... -0.67483383 -0.68801788
  1.31667716]
 [ 0.05003309  0.12213032 -0.10799087 ...  1.48184628 -0.68801788
 -0.75948762]
 ...
 [ 0.6329849   0.67640943  1.56576738 ... -0.67483383 -0.68801788
  1.31667716]
 [ 0.84118198  0.77718745  0.25067161 ...  1.48184628 -0.68801788
 -0.75948762]
 [ 1.54905203  1.48263359  1.32665906 ... -0.67483383 -0.68801788
  1.31667716]]
```

10. Split the data into training and testing

In [321]:

```
X.shape, y.shape
```

Out[321]:

```
((4177, 10), (4177,))
```

In [322]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X,y, test_size=0.2, random_state=42)
```

In [323]:

```
print(' x_tain.shape : ',x_train.shape)
print(' y_tain.shape : ',y_train.shape)
print(' x_test.shape : ',x_test.shape)
print(' y_test.shape : ',y_test.shape)

x_tain.shape : (3341, 10)
y_tain.shape : (3341,)
x_test.shape : (836, 10)
y_test.shape : (836,)
```

10. Build the Model, 11. Train the Model , 12.Test the Model

In [324]:

```
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(x_train, y_train)
lr_pred = lr.predict(x_test)
```

In [325]:

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error,make_scorer
from sklearn.model_selection import RandomizedSearchCV

rf = RandomForestRegressor()

param = {
    'max_depth':[3,6,9,12,15],
    'n_estimators':[10,50,100,150,200]
}
rf_search = RandomizedSearchCV(rf,param_distributions=param,n_iter=5,scoring=make_scorer
(mean_squared_error),n_jobs=-1,cv=5,verbose=3)
rf_search.fit(x_train, y_train)
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

Out[325]:

```
RandomizedSearchCV(cv=5, estimator=RandomForestRegressor(), n_iter=5, n_jobs=-1,
    param_distributions={'max_depth': [3, 6, 9, 12, 15],
    'n_estimators': [10, 50, 100, 150, 200]},
    scoring=make_scorer(mean_squared_error), verbose=3)
```

In [326]:

```
means = rf_search.cv_results_['mean_test_score']
params = rf_search.cv_results_['params']
for mean, param in zip(means, params):
    print("%f with: %r" % (mean,param))
    if mean == min(means):
        print('Best parameters with the minimum Mean Square Error are:' ,param)

4.664623 with: {'n_estimators': 200, 'max_depth': 6}
4.618707 with: {'n_estimators': 100, 'max_depth': 15}
4.644619 with: {'n_estimators': 200, 'max_depth': 15}
5.677870 with: {'n_estimators': 150, 'max_depth': 3}
4.581780 with: {'n_estimators': 100, 'max_depth': 9}
Best parameters with the minimum Mean Square Error are: {'n_estimators': 100, 'max_depth': 9}
```

In [327]:

```
rf = RandomForestRegressor(n_estimators=50,max_depth=6)
rf.fit(x_train,y_train)

rf_pred = rf.predict(x_test)
```

14. Measure the performance using Metrics

In [328]:

```
from sklearn import metrics
print('Linear Regression :')
print('.....')
print('MAE:',metrics.mean_absolute_error(y_test, lr_pred))
print('MSE:',metrics.mean_squared_error(y_test, lr_pred))
print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test, lr_pred)))
print('R2 Score:',metrics.r2_score(y_test,lr_pred))
print('\n\n')
```

Linear Regression :

```
.....
MAE: 1.5944508821770336
MSE: 4.892375672262822
RMSE: 2.211871531591024
R2 Score: 0.5480572061259404
```

In [329]:

```
from sklearn import metrics
print('Random Forest Contains:')
print('.....')
print('MAE:',metrics.mean_absolute_error(y_test, rf_pred))
print('MSE:',metrics.mean_squared_error(y_test, rf_pred))
print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test, rf_pred)))
print('R2 Score:',metrics.r2_score(y_test,rf_pred))
```

Random Forest Contains:

```
.....
MAE: 1.5580369509719958
MSE: 5.025592967383406
RMSE: 2.241783434541215
R2 Score: 0.535750997326301
```